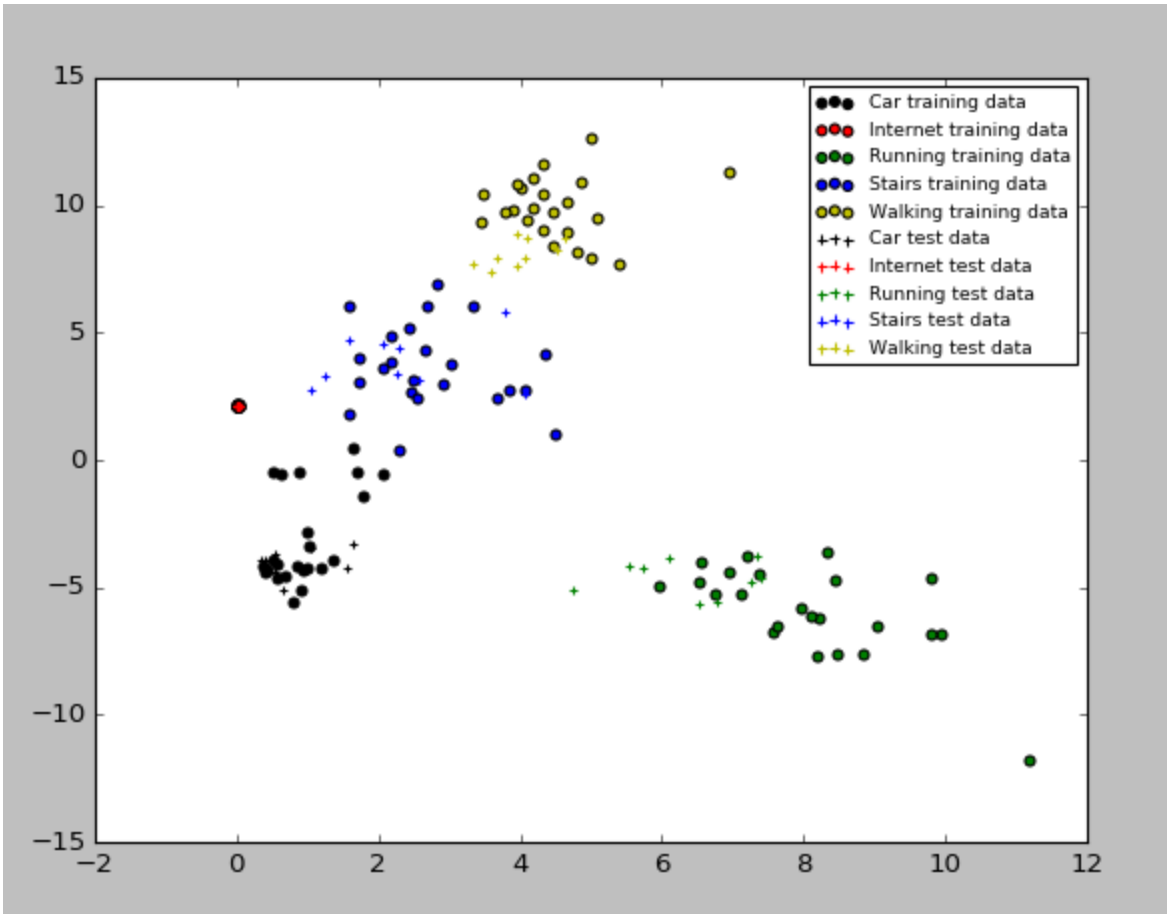

Logistic Regression

This algorithm got 49/50 predictions right. I incorrectly classified one of the “running” test data points as a “car” data point.

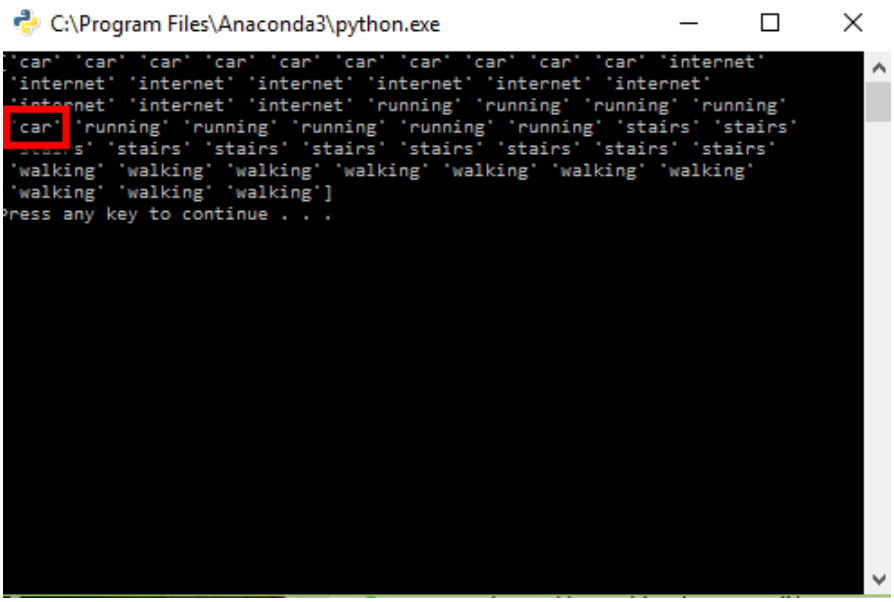

K Neighbors and SVM

Both of these algorithms successfully classified 50/50 test points. I think this worked better for this data in particular because there are clear clusters in the data—the boundary lines are not ambiguous.

Linear regression would probably work better in a case where the data is not in such clear clusters, and a regression line needs to be found that actually fits the data.